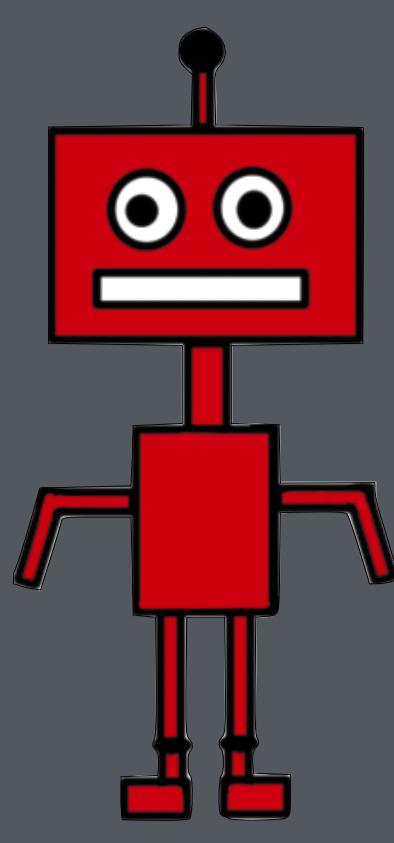


Learning Dirichlet Priors for Affordance Aware Planning

David Abel & Gabriel Barth-Maron, James MacGlashan, Stefanie Tellex
Dept. of Computer Science, Brown University

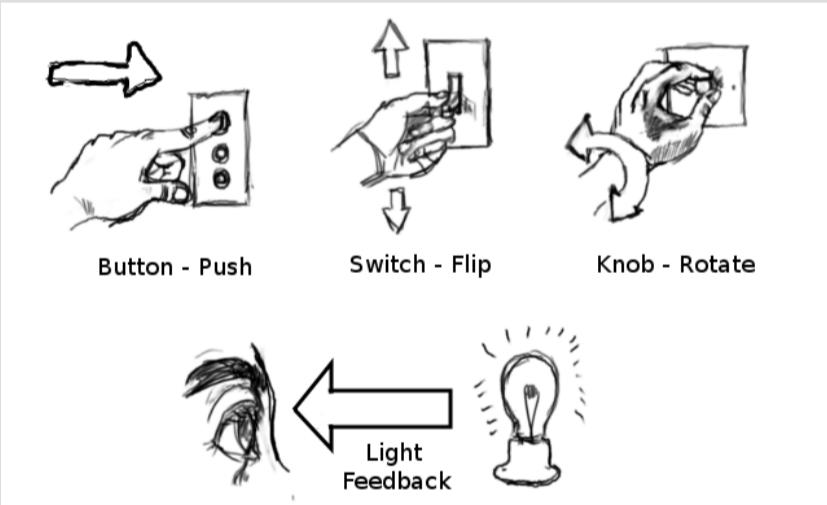


Goal

Enable autonomous agents to learn how to plan efficiently in massive stochastic state spaces.

Background

Affordances: Direct agent toward relevant action possibilities.



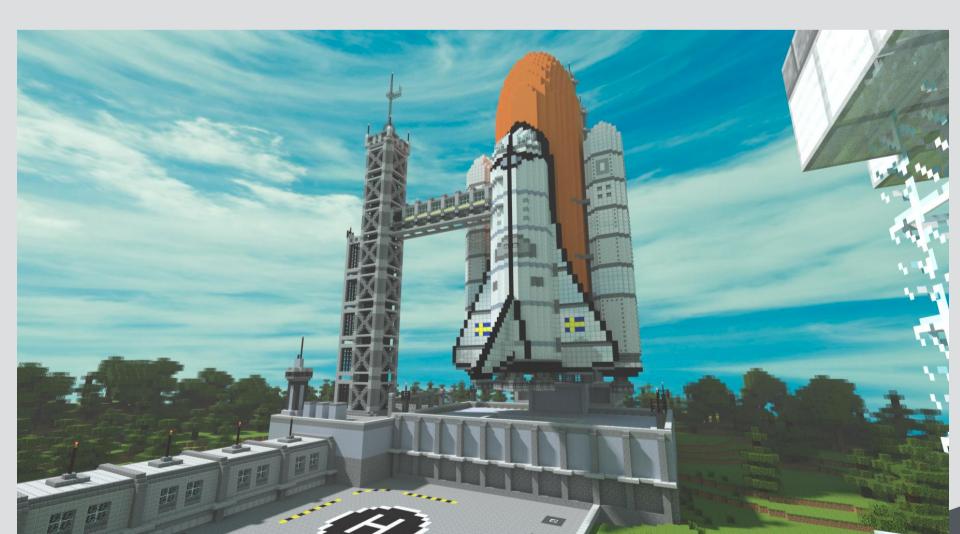
"What [the environment] offers [an] animal, what [the environment] provides or furnishes, either for good or ill"
- J.J. Gibson, 1977

Formalism:

$$\Delta = \langle p, g \rangle \longmapsto \mathcal{A}'$$

p = predicate on states
 g = lifted goal description
 \mathcal{A}' = subset of MDP Actions

Domain: Minecraft



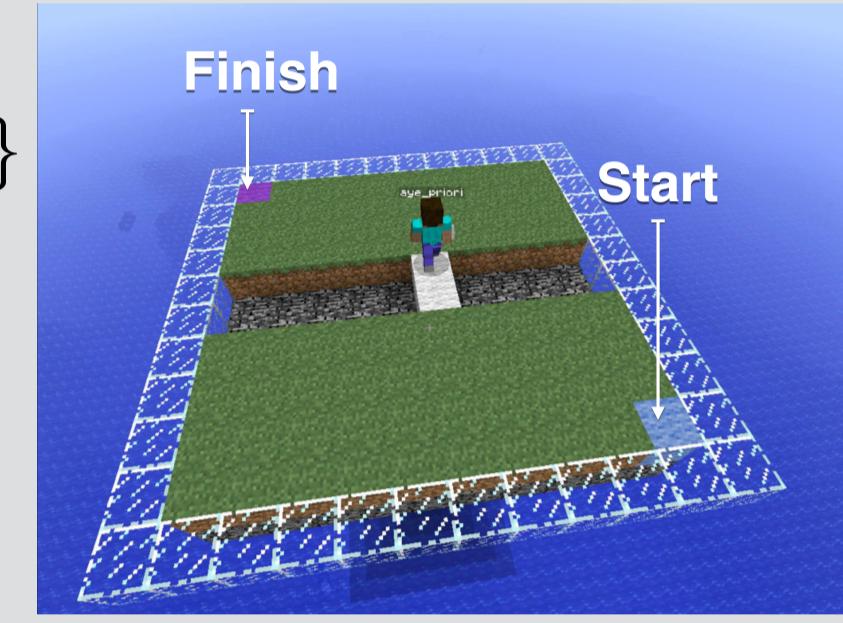
\approx Turing Complete Legos

Affordance Example

$$\Delta_1 = \langle \text{nearPlane}, \text{atLoc} \rangle \longmapsto \{\text{move}\}$$

$$\Delta_2 = \langle \text{nearTrench}, \text{atLoc} \rangle \longmapsto \{\text{place}\}$$

If Δ 's predicate is true and Δ 's goal type matches the current goal, use Δ 's actions

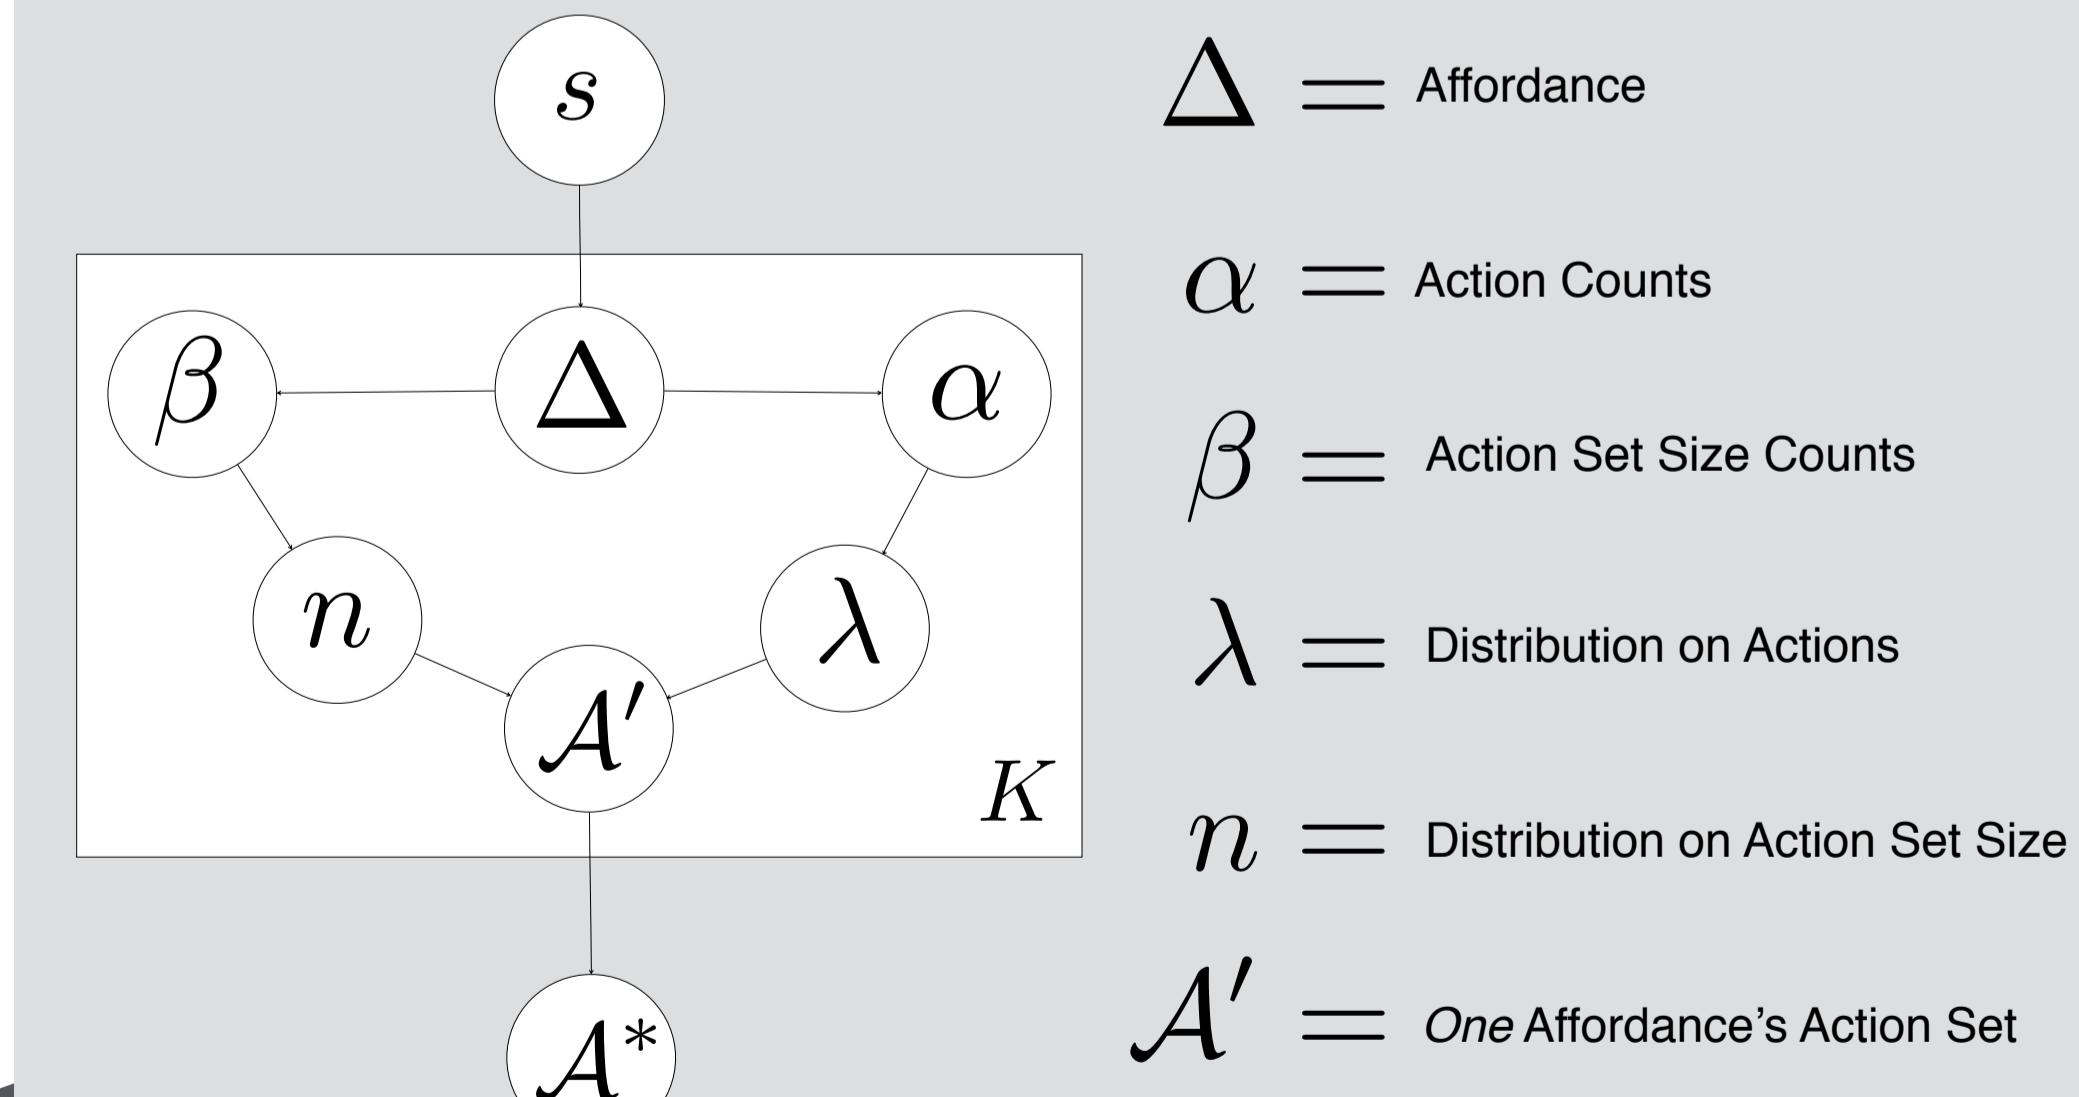


Learning

Goal: For a given state, for each affordance, learn which actions are most relevant:

$$\Pr(\mathcal{A}^* \mid s, \Delta_1 \dots \Delta_K)$$

Graphical Model:



$$\Pr(\lambda \mid \alpha) = \text{DirMult}(\alpha)$$

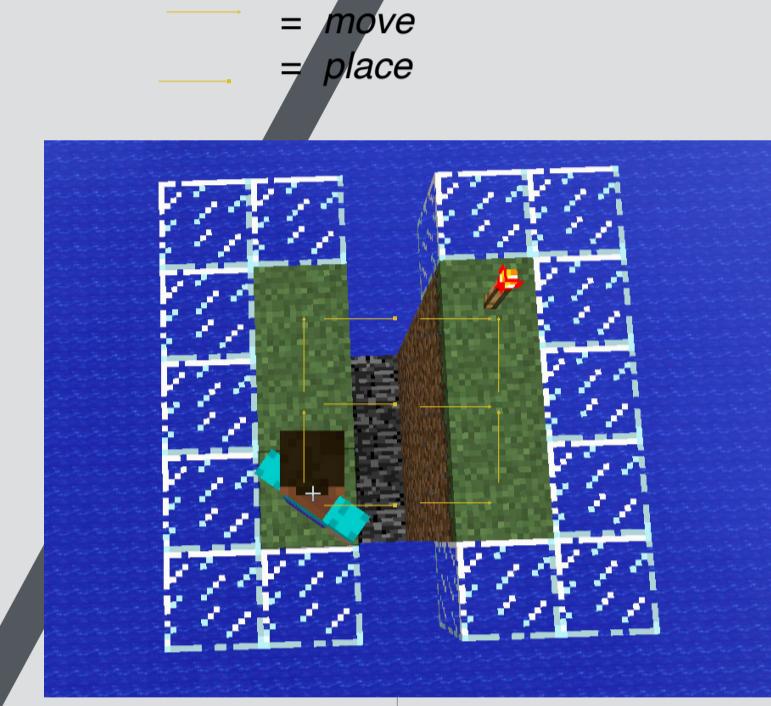
$$\Pr(n \mid \beta) = \text{Dir}(\beta)$$

Learning Example

1) For each activated affordance, count:

α = number of worlds in which each action was used

β = number of unique actions used in each world



$\Delta_1 = \langle \checkmark \text{ nearTrench}, \checkmark \text{ atGoal} \rangle$

$$\Delta_1.\alpha.\text{moveRight}++, \Delta_1.\alpha.\text{moveForward}++, \Delta_1.\alpha.\text{placeRight}++$$

$$\Delta_1.\beta.3++$$

2) When solving the MDP on a new state space, in each state s :

$$\mathcal{A}^* = \bigcup_{i=1}^K (\Delta_i.\text{getActions}(s))$$

3) Where

$$\Delta_i.\text{getActions}(s):$$

```

 $\lambda \leftarrow \text{DirMult}(\Delta_i.\alpha)$ 
 $n \leftarrow \text{Dir}(\Delta_i.\beta)$ 
 $\mathcal{A}' \leftarrow_n \lambda$ 
return:  $\mathcal{A}'$ 

```

Results

